

Application Life Cycle Management

Rolling Update and Rollback - Part 2

Doing this way will result in the deployment definition file a different configuration.

The difference between the recreate and rolling-update strategies can also be seen when the deployments are viewed in detail.

"unbeatable describe deployment 'deployment name'"
When a deployment is created, the old deployment is scaled down to '0'. See in the ~~the~~ message the column of the 'event' section ~~and~~ of the command mentioned above. Then the replica set scale up to the mentioned number.

However, when the rolling update is used, the old replica set was scaled down one at a time, simultaneously scaling up the new replica set one at a time.

(2)

07/07/24

Upgrade

When a deployment is created, say to deploy five replicas, it first creates a replica set automatically, which in turn creates the number of pods required to meet the number of replicas. When the app is upgraded, the Kubernetes deployment object creates a new replica-set under the hood and starts deploying the containers there. At the same time, taking down the pods in the old replica set, following a rolling update strategy. This can be seen ~~when~~ using the following command -

```
kubectl get replicsets
```

The old replica-set with 0 pods and the new replica-set with 5 pods.

Kubernetes deployment allows to rollback to a previous version, revision. To undo a change

(3)

07/09/24

kubectl rollout undo 'deployment name'

The deployment will then destroy the pods in the new replica set, and bring the older ones up in the old replica-set.

Again "kubectl get deployments 'deployment name'" will show the old deployment has now the desired number of pods.

Rolling Update ~~conf~~

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: my-app
```

spec:

replicas: 3

strategy:

type: RollingUpdate

rollingUpdate:

maxUnavailable: 1

maxSurge: 1

(Same level as strategy)

selector:

matchLabels:

app: my-app

template:

metadata:

labels:

app: my-app

spec:

containers:

- name: my-container

image: my-image:latest

ports:

- containerPort: 80

⑨

07/07/24

~~Rolling~~ Recreate

kind: Deployment

specVersion: apps/v1

kind: Deployment

metadata:

name: my-app

spec:

replicas: 3

strategy:

type: Recreate

selector:

matchLabels:

app: my-app

template:

metadata:

labels:

app: my-app

spec:

containers:

- name: my-container

image: my-image:latest

ports:

- containerPort: 80

Commands and Arguments in Kubernetes

The dockerfile has 'ENTRYPOINT' and 'CMD' command. The entrypoint is the command that is run at startup, and the 'CMD' is the default parameter passed the command. With the 'args' option in the ~~pod~~ pod definition file, it overrides the 'CMD' instruction in the dockerfile.

To override the entrypoint command use 'command' option in the pod definition file. The 'command' field corresponds to entrypoint instruction in the dockerfile.

Ubuntu dockerfile

FROM ubuntu

ENTRYPOINT ~~["sleep"]~~ ["sleep"]

CMD ["5"]

pod-definition.yaml

apiVersion: v1

kind: Pod

metadata:

name: ubuntu-sleeper-pod

spec:

containers:

- name: ubuntu-sleeper

image: ubuntu-sleep

command: ["sleep 20"]

args: ["10"]

07/07/29

⑥

apiVersion: v1

kind: Pod

metadata:

name: custom-command-pod

spec:

containers:

- name: my-container

image: Ubuntu

command: ["sh", "-c"]

args: ["echo Hello, Kubernetes! && sleep 3600"]

Environment variables in Kubernetes

pod-definition.yaml

apiVersion: v1

kind: Pod

metadata:

name: simple-webapp-color

spec:

containers:

- name: ~~ubuntu~~ simple-webapp-color

image: simple-webapp-color (Same level as 'ports')

ports:

- containerPort: 8080

env:

- name: APP_COLOR

value: pink

(7)

07/07/24

Config Maps

When there is a lot of pod definition files it becomes difficult to manage the environment data stored within the query files. These information can be taken out of the pod definition files and manage it centrally using Configuration Maps.

Configmaps are used to pass configuration data in the form of key-value pairs in Kubernetes.

```
kubectl create configmap \ app-config --from-literal=APP-COLOR=blue
```

```
kubectl create configmap <config-name> --from-file=path-to-file>
```

config-map.yaml

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: app-config
```

```
data:
  APP-COLOR: blue
  APP-MODE: prod
```

(8)

07/07/24

Use different config maps for different apps.

kubectl get configmaps

pod-definition.yaml

apiVersion: v1

kind: Pod

metadata:

name: simple-webapp-color

labels:

name: simple-webapp-color

spec:

containers:

- name: simple-webapp-color

image: simple-webapp-color

ports:

- containerPort: 8080

envFrom:

- configMapRef:

name: app-config

Secrets

Secrets are used to store sensitive information like passwords or keys. They are stored in a encrypted or hashed format. As with configMap, there are two steps involved in working with secrets. First, create the secret and second inject them into pod.

There are two ways to create a secret - the imperative way without using a secret definition file and the declarative way by using a secret definition file.

kubectl create secret generic <secret-name> --from
--from-literal=<key>=<value>

kubectl create secret generic & app-secret
--from-literal=DB_HOST=mysql

For additional key-value pairs, use '--from-literal' multiple times. However this could get complicated when there is too many secrets to pass in. Another

(10)

06/07/24

way to input the secret data is through a file.

Run `kubectl create secret generic app-secret --from-file=app-secret.properties`

For the declarative approach, create a definition file.

`secret-data.yaml`

```
apiVersion: v1
kind: Secret
metadata:
  name: app-secret
data:
  DB_HOST: mysql
  DB_USER: root
  DB_PASSWORD: password
```

But the values are in plain text, so they need to be in encoded format.

To convert into encoded format

```
[echo -n 'mysql' | base64]
```

Output: NjklZcWw =

kubectl get secrets

kubectl describe secrets

To see the secrets with their values

kubectl get secret app-secret -o yaml

To decode the Base64 encode

echo -n 'BxLzCWw=' | base64 --decode

Output: mysql

Inject the encoded secret into the pod definition file

pod-definition.yaml

apiVersion: v1

kind: Pod

metadata:

name: simple-webapp-color

labels:

name: simple-webapp-color

spec:

containers:

- name: simple-webapp-color

image: simple-webapp-color

ports:

- containerPort: 8080

(Same level as envFrom

- envFrom

- secretRef:

name: app-secret

All the secret files can be injected in a single environment variables or inject the whole secret as files in a volume secrets in a volume.

apiVersion: v1

kind: Pod

metadata:

name: secret-volume-pod

spec:

containers:

- name: mycontainer

image: nginx

volumeMounts:

- name: mysecretvolume

mountPath: /etc/secret-volume

volumes:

- name: mysecretvolume

secret:

: secretName: mysecret

Note on secrets

- secrets are not encrypted. Only encoded.
- Don't check-in secret objects to SCM along with code.
- secrets are not encrypted in 'etcd'.
 - ↳ Enable encryption at rest
- Anyone able to create pods/deployments in the same namespace can access the secrets.
 - ↳ configure least-privilege access to Secrets-RBAC
- consider third-party secrets store providers
 - AWS provider, Azure provider, GCP provider, vault provider.

08/07/24Additional Resource

youtube.com/watch?v=MTn9W9MxnRI

- The way Kubernetes handles secrets, such as:
 - A secret is only sent to a node if a pod on that node requires it.
 - Kubernetes stores the secret into a tmpfs so that the secret is not written to disk storage
 - Once the pod that depends on the secret is deleted,

Kubelet will delete its local copy of the secret data as well.

There are other ways of handling sensitive data like passwords in Kubernetes, such as using tools like Helm secrets, and HashiCorp Vault.

Multi-Micro-container Pods

The idea of decoupling a large monolithic application into sub-components known as microservices enables to develop and deploy a set of independent, small, and reusable code. This architecture can then help to scale up, down as well as modify each service as required - as opposed to modifying the entire app.

However, at times it may be necessary for two services to work together, such as a web server and a logging service. One agent instance per webserver instance needs to be paired together. Don't merge and bloat the code of the two services,

as each of them target different functionalities. The requirement here is the two functionality to work together.

For these requirements, multi-container pods are necessary that share the same life cycle. These share the same network space which means they can refer to each other as localhost, and they have access to the same storage volumes.

This way, no need to establish volume sharing or services between the pods to enable communication between them.

pod-definition.yaml

apiVersion: v1

kind: Pod

metadata:

name: simple-webapp

labels:

name: simple-webapp

image: simple-webapp

ports:

- containerPort: 8080

spec:

containers:

- name: simple-webapp

image: simple-webapp

ports:

- containerPort: 8080

- name: log-agent

image: log-agent

Init containers

In a multi-container pod, each container is expected to run a process that stays alive as long as the pod's lifecycle.

For example, in the multi-container pod, there is a web app and logging agent, both the containers are expected to stay alive at all times.

The process running in the log agent container is expected to stay alive as long as the web-application is running. If any of them fail, the pod restarts. But, a process runs in a container that must be completed. For example, a process that pulls a code or library from a repository that will be used by the main web app.

That is a task that will be run only one time when the pod is first created, or a process that waits for an external service or db to be up before the actual app starts.

That's where `initContainer` is configured in a pod like all other containers, except that is specified inside a `'initContainers'` section.

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: myapp-pod
```

```
  labels:
```

```
    app: myapp
```

```
spec:
```

```
  containers:
```

- name: myapp-container

- image: busybox:1.28

- command: [sh, '-c', 'echo The app is running']

```
  initContainers:
```

- name: init-myservice

- image: busybox

- command: [sh, '-c', 'git clone ;']

When a pod is first created the `initContainer` is run, and the process in the `initContainer` must run to a completion before the real container

08/07/24

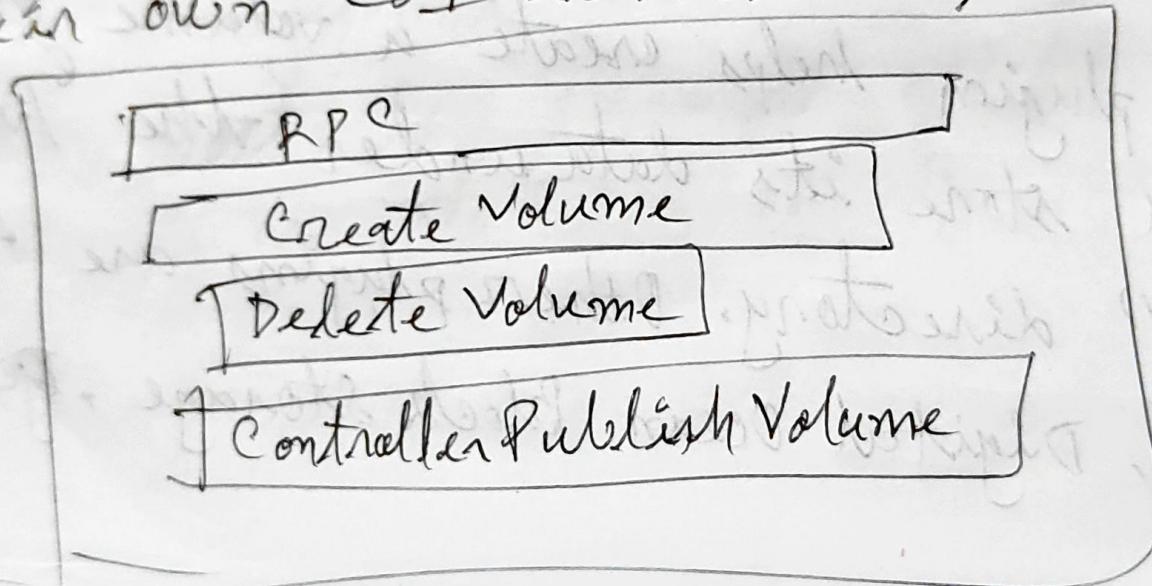
(18)

container hosting the app starts.

Such ~~for~~ multiple 'initContainers' can be configured. In that case, each init container is run one at a time in sequential order. If any of the init containers fail to complete, Kubernetes restarts the pod repeatedly until the init-container succeeds.

Container storage Interface

The container runtime interface (CRI) is a standard that defines how an orchestration solution like Kubernetes would communicate with container runtimes like docker. To extend support for different docker networking solution, the CNI was introduced. Every vendor has their own CSI driver. \nearrow CSI procedure



Volumes

The pods created in Kubernetes are transient in nature. When a ~~pod~~ pod is created to process data and then deleted, the data processed by it gets deleted as well. For this, a volume to the pod is attached. The data generated by the pod is now stored in the volume, even after the pod is deleted, the data remains.

There is a single node Kubernetes cluster, create a pod that generates a random between 1 and 100 and writes that to a file, then gets deleted along with the random number.

```

apiVersion: v1
kind: Pod
metadata:
  name: random-number-generator
spec:
  containers:
    - image: alpine
      name: alpine
      command: ["/bin/sh", "-c"]
      args: ["shuf -i 0-100 -n 1 >/opt/number.out;"]
  volumeMounts:
    - mountPath: /opt
      name: data-volume

```

To retain the number generated by the pod - create a volume and a volume needs storage. Once the volume is created, to access it from a container, the volume is mounted to a directory inside the container.

Use the volume's 'volumeMounts' field in each container to mount the 'data-volume' to the directory '/opt' within the container. The random number will now be written to '/opt' mounted inside the container which happens to be 'data-volume' which is in fact '/data' directory on the host. When the pod gets deleted, the file with the random number still lives on the host.

(Same level as containers)
volumes:

- name: data-volume

hostpath:

path: /data

type: Directory

rights: appnd
rights: rmw

rights: rwm

rights: rwm

rights: rwm

rights: rwm

The way the volume is configured here is not recommended for use in a multimode cluster. This is because the pods would use the '/data' directory on all the nodes, and expect all of them to be the same and have the same data since they are on different servers. They are ~~not~~ in fact not the same unless configured some kind of replicated external replicated cluster storage solution. Kubernetes supports several types of standard storage solutions such as NFS, glusterFS, cephFS, AWS EBS, Azure Disk etc.

09/07/24

Persistent Volumes

~~the~~ As per the previous method, in a large environment with a lot of users deploying a lot of pods the users would have to configure storage everytime for each pod. Everytime, any change to be made, the users would have to make them on all of his pods. Persistent volume provides centralized solution.

A persistent volume is a cluster wide pool of storage volumes configured by an administrator

to be used by users deploying apps on the cluster. The users can now select storage from this pool using persistent volume claims.

PV-definition.yaml

```
apiVersion: v1
```

```
kind: PersistentVolume
```

```
metadata:
```

```
  name: pv-val1
```

```
spec:
```

```
  accessModes:
```

- ReadWriteOnce

```
  capacity:
```

```
    storage: 1Gi
```

```
    hostPath:
```

```
      path: /tmp/data
```

The supported 'accessModes' are 'ReadOnlyMany', 'ReadWriteOnce' and 'ReadWriteMany'.

kubectl create -f PV-definition.yaml

kubectl get persistentvolume

Persistent Volume Claims

Persistent volume and persistent volume claims are two separate objects in Kubernetes namespace. An admin creates a set of persistent volumes and a user creates persistent volume claims to use storage. Once the persistent volume claims are created, Kubernetes binds the persistent volume to 'claims' based on the request and properties set on the volume. Every persistent volume claim is bound to single persistent volume. During the binding process, Kubernetes tries to find a persistent volume that has sufficient capacity as requested by the 'claim' and any other properties such as access modes, volume modes, storage class etc. However, if there are multiple possible matches for a single claim and I would like to specifically use a particular volume, labels and selectors still could be used to bind to the right volumes. A smaller claim may get bound to a larger volume if all other

criteria matches and there are no better options. There is one-to-one relationship between claims and volumes, so no other claims can utilize the remaining capacity in the volume.

If there are no volumes available, the persistent volume claim will remain in a pending state until newer volumes are made available to the cluster once newer volumes are available. Once newer volumes are available, the claim would automatically be bound to the newly available volume.

pvc-definition.yaml

apiVersion: v1

kind: PersistentVolumeClaim

metadata:

name: myclaim

spec:

accessModes:

- ReadWriteOnce

resources:

request:

storage: 500 Mi

kubectl apply -f pvc-definition.yaml

kubectl get persistentvolumeclaim

When the claim is created, Kubernetes looks at the volume created previously. The 'accessMode' matches. The capacity requested is 500 MB but the volume is configured with 1 GB of storage. Since there are no other volumes available, the persistent volume claim is bound to persistent volume.

Delete PVCS

kubectl delete persistentvolumeclaim myclaim

The admin can choose what will happen when the 'claim' is deleted: By default it is set to retain meaning the persistent volume will remain until it is manually deleted by the admin. It is not available for reuse by any other 'claim' or it can be deleted automatically. This way as soon as the claim is deleted, the volume will be deleted as well, thus freeing up storage on the end storage.

device, as a third option to recycle. In this case, in the data volume will be scrubbed before making it available to other claims. one in pod definition file

apiVersion: v1

kind: Pod

metadata:

name: mypod

spec:

containers

- name: myfrontend

image: nginx

VolumeMounts:

- mountPath: "/var/www/html"

name: mypd

volumes:

- name: mypd

persistentVolumeClaim:

claimName: myclaim

deployment definition file

```
apiVersion: apps/v1
```

```
kind: Deployment
```

```
metadata:
```

```
  name: mydeployment
```

```
spec:
```

```
  replicas: 3
```

```
  selector:
```

```
    matchLabels:
```

```
      app: myapp
```

```
  template:
```

```
    metadata:
```

```
      labels:
```

```
        app: myapp
```

```
  spec:
```

```
    containers:
```

```
      - name: myfrontend
```

```
        image: nginx
```

```
    volumeMounts:
```

```
      - mountPath: "/var/www/html"
```

```
        name: mypd
```

```
    volumes:
```

```
      - name: mypd
```

```
    persistentVolumeClaim:
```

```
      claimName: myclaim
```

(20)

09/07/24

Storage Class

With storage classes, provisioner can be defined that can automatically provision storage and attach them to a pod when a claim is made. That's called dynamic provisioning of volumes. Create a storage class definition file, the pod is in the ~~go~~ google cloud

sc-definition.yaml

apiVersion: storage.k8s.io/v1

kind: StorageClass

metadata:

name: google-storage

provisioner: kubernetes.io/gce-pod

(same level as provisioner)

parameters:

type: pd-standard

replication-type:

none

So now the PV definition is not needed, because the PV and any associated storage is going to be created automatically when the storage class is created. The PV definition file -

(39)

09/07/24

pvc - definition, yaml

apiVersion: v1

kind: PersistentVolumeClaim

metadata:

name: myclaim

spec:

accessModes:

- ReadWriteOnce

storageClassName: google-storage

resources:

requests:

storage: 500Mi

That's how the pvc knows which storage class to use. Next time, if pvc is created, the storage class associated with it uses the defined provisioner to provision a new disk with the required size. Then creates a persistent volume and then binds the pvc to that volume. With different types of provisioner, pass different types of parameters.